


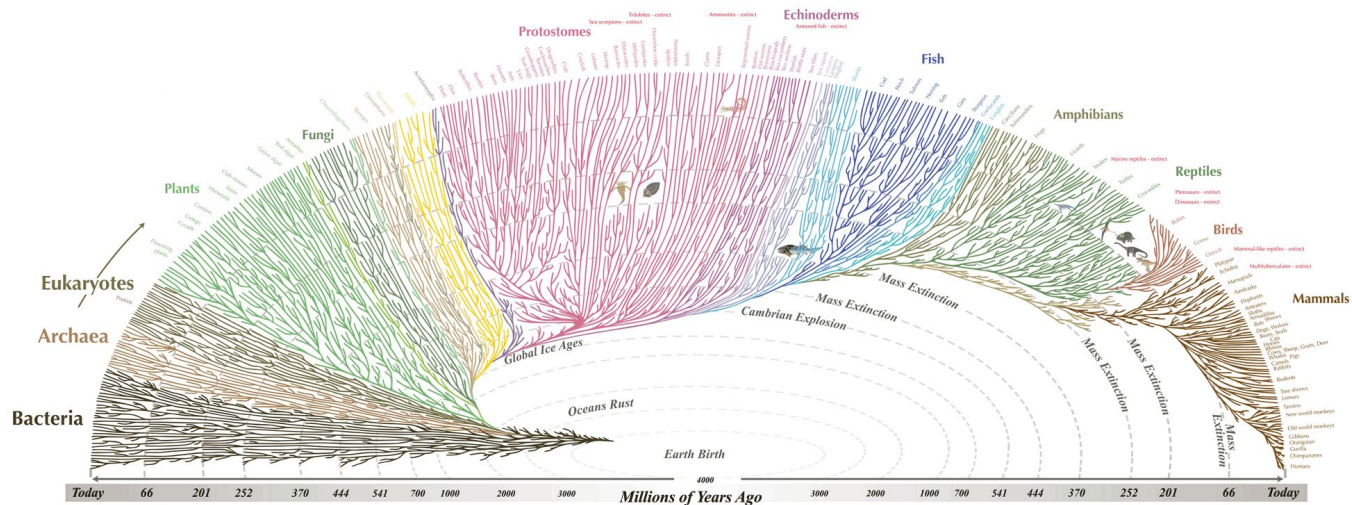
# Week 4: Sequence alignment

- Global alignment
  - Dynamic programming
  - Needleman-Wunsch algorithm
- Local alignment
  - Smith-Waterman algorithm
  - BLAST

# Sequence evolution



Then between A & B. various  
 loss of relation. C & B. the  
 first predation, B & D  
 rather greater distance  
 then former would have  
 formed. - heavy relation



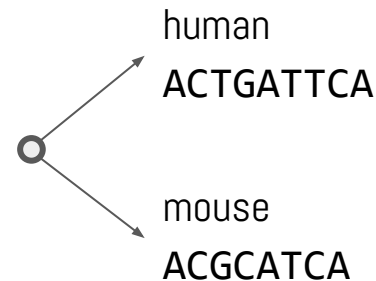
All the major and many of the minor living branches of life are shown on this diagram, but only a few of those that have gone extinct are shown. Example: Dinosaurs - extinct

© 2008, 2017 Leonard Eisenberg. All rights reserved.  
 eisenberg.com


deletion      mutation      insertion

ACATGGTCA → AC\*TGGTCA → ACTGATCA → ACTGATTCA

Evolutionary time



# Sequence alignment



Sequences can be aligned by allowing for **gaps** and **mismatches**.

ACTGATTCA

ACGCA-TCA

ACTGATTCA

AC-GCATCA

ACTG-ATTCA

AC-GCAT-CA

Which alignment is correct?

A scoring scheme:

- Match: 2
- Mismatch: -3
- Gap: -2

*We will come back to this!*

2+2-3-3+2-2+2+2+2

= 4

2+2-2+2-3-3+2+2+2

= 4

2+2-2+2-2+2+2-2+2+2

= 8

**Alignment is gap placement.**

How many possible alignments?

# Dynamic programming

Solve a given complex problem by:

1. Breaking it into subproblems and
2. Storing the results of subproblems to avoid computing the same results again.

Two key properties of a problem that suggest that the given problem can be solved using DP.

1. Overlapping Subproblems
  - Given problem can be recursively broken down into subproblems that can be related to each other. That is, total no. of subproblems is polynomial.
2. Optimal Substructure
  - The optimal solution can be produced by combining optimal solutions of subproblems.



Richard Bellman

Optimal decision processes, involved time series & planning - thus 'dynamic' & 'programming'.

"It's impossible to use the word dynamic in a pejorative sense"; DP was "something not even a Congressman could object to."


# Dynamic programming

Hemachandra/Fibonacci numbers: 0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144, .....

$$\begin{aligned} F_0 &:= 0; F_1 := 1; \\ F_n &= F_{n-1} + F_{n-2}, \text{ for all } n \geq 2. \end{aligned}$$

A trivial algorithm for computing  $F_n$ :

```
naive_fib(n):  
    if n ≤ 1: return n  
    else: return naive_fib(n - 1) +  
           naive_fib(n - 2)
```



# Dynamic programming


Hemachandra/Fibonacci numbers:  $F_0 := 0$ ;  $F_1 := 1$ ;  $F_n = F_{n-1} + F_{n-2}$ , for all  $n \geq 2$ .

Never recompute a subproblem  $F(k)$ ,  $k \leq n$ , if it has been computed before.

Memoization: Remembering previously computed values.

Improved algorithm for computing  $F_n$ :

```
memo = { }  
  
fib(n):  
    if n in memo: return memo[n]  
    else if n = 0: return 0  
    else if n = 1: return 1  
    else: f = fib(n - 1) + fib(n - 2)  
    memo[n] = f  
    return f
```



These values are already computed and stored in memo when runtime processes these nodes of the recursion.

# Dynamic programming

1. Overlapping Subproblems
2. Optimal Substructure

DP  $\approx$  recursion + memoization (reuse)

- Remember (memoize) previously solved “subproblems”; e.g., in Fibonacci, we memoized the solutions to the subproblems  $F_0, F_1, \dots, F_{n-1}$ , while unraveling the recursion.
- If we encounter a subproblem that has already been solved, reuse solution.
- Runtime  $\approx$  (no. of subproblems) \* (time per subproblem)

# Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

Align **GCAT** with **GAT**

## Step 1

A scoring scheme:

- Match: 1
- Mismatch: -2
- Gap: -1

	—	G	C	A	T
—					
G					
A					
T					



# Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

Align GCAT with GAT

$$M(0, j) = j * p$$

Step 2

$$M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left( \begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + S(A_i, B_j) \end{array} \right)$$

top  
left  
diagonal

	—	G	C	A	T
—					
G					
A					
T					

# Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

Align GCAT with GAT

$$M(0, j) = j * p$$

Step 2

$$M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left( \begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + S(A_i, B_j) \end{array} \right)$$

top  
left  
diagonal

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1				
A	-2				
T	-3				

# Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

$$M(0, j) = j * p$$

Step 2

$$M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left( \begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + S(A_i, B_j) \end{array} \right)$$

top

left

diagonal

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1	?			
A	-2				
T	-3				

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1	-2			
A	-2				
T	-3				

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1	-2			
A	-2				
T	-3				

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1	1			
A	-2				
T	-3				

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1	1			
A	-2				
T	-3				

# Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

Align GCAT with GAT

$$M(0, j) = j * p$$

Step 2

$$M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left( \begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + S(A_i, B_j) \end{array} \right)$$

top

left

diagonal

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1	1	0	-1	-2
A	-2	0	0	1	0
T	-3	-1	-2	0	2

# Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

Align GCAT with GAT

GCAT  
G-AT

$$M(0, j) = j * p$$

$$M(i, 0) = i * p$$

$$M(i, j) = \text{MAX} \left( \begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + S(A_i, B_j) \end{array} \right)$$

top

left

diagonal

Step 3

	—	G	C	A	T
—	0	-1	-2	-3	-4
G	-1	1	0	-1	-2
A	-2	0	0	1	0
T	-3	-1	-2	0	2

# Needleman-Wunsch algorithm

1. Scoring function: substitution matrix & gap penalty
2. Matrix initialization & filling
3. Traceback

Align **ATGCT** with **ATTACA**

$$M(0, j) = j * p$$

$$M(i, 0) = i * p$$


$$M(i, j) = \text{MAX} \left( \begin{array}{l} M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + S(A_i, B_j) \end{array} \right)$$

top  
left  
diagonal


	-	A	T	T	A	C	A
-							
A							
T							
G							
C							
T							

# Global & local alignment


A local alignment of strings  $s$  and  $t$  is an alignment of a substring of  $s$  with a substring of  $t$ .



Global Alignment



Local Alignment



# Smith-Waterman algorithm

Similar to Needleman-Wunsch, with 3 changes:

- First row/column set to 0.
- No negative scores, set to 0.
- Backtrack from cell with highest score, stop at 0.

Align GCAT with GCT

$$M(0, j) = 0$$

$$M(i, 0) = 0$$

$$M(i, j) = \text{MAX} \left( \begin{array}{l} 0, \\ M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + S(A_i, B_j) \end{array} \right)$$

top  
left  
diagonal

	-	G	C	A	T
-					
G					
C					
T					



# Smith-Waterman algorithm

Similar to Needleman-Wunsch, with 3 changes:

- First row/column set to 0.
- No negative scores, set to 0.
- Backtrack from cell with highest score, stop at 0.

$$M(0, j) = 0$$

$$M(i, 0) = 0$$

$$M(i, j) = \text{MAX} \left( \begin{array}{l} 0, \\ M(i-1, j) + p, \\ M(i, j-1) + p, \\ M(i-1, j-1) + S(A_i, B_j) \end{array} \right)$$

top

left

diagonal

Align GCAT with GCT

GC  
GC

	-	G	C	A	T
-	0	0	0	0	0
G	0	1	0	0	0
C	0	0	2	1	0
T	0	0	1	1	2

# Substitution matrix to measure similarity in sequence alignments



Margaret Dayhoff

Applying math & computational techniques to the sequencing of proteins and nucleic acids.

- 1965: First collection of protein seqs.
- Single-letter code for amino acids.
- 1966: 'Evolutionary trees'.
- **1978: First AA similarity-scoring matrix.**
- 1980: Launched the Protein Information Resource, the first online database system that could be accessed by telephone line.

**Substitution matrix:** A collection of scores for aligning nucleotides or amino acids with one another.

- The scores represent the relative ease with which one nucleotide or amino acid may mutate into or substitute for another.
- Purely statistical, nothing directly to do with structure/biochemistry.

[illegible]

# Substitution matrix to measure similarity in sequence alignments

**Substitution matrix:** Each scores is a log-odds score equal to the logarithm of the ratio of the likelihoods of two hypotheses: the residues can substitute for one another or not.

$$s(a,b) = \frac{1}{\lambda} \log \frac{p_{ab}}{f_a f_b}$$

- $p_{ab}$ : likelihood of these two residues being correlated because they're homologous.
  - $p_{ab}$  are the target frequencies: the probability that we expect to observe residues  $a$  and  $b$  aligned in homologous sequence alignments.
- $f_a f_b$ : likelihood of these two residues being uncorrelated and unrelated, occurring independently.
  - $f_a$  and  $f_b$  are background frequencies: the probabilities that we expect to observe amino acids  $a$  and  $b$  on average in any protein sequence.
- $\lambda$ : a scaling factor, usually set to something that lets helps round off all the terms in the score matrix to sensible integers.

# Substitution matrix to measure similarity in sequence alignments

**Substitution matrix:** Each scores is a log-odds score equal to the logarithm of the ratio of the likelihoods of two hypotheses: the residues can substitute for one another or not.

$$s(a,b) = \frac{1}{\lambda} \log \frac{p_{ab}}{f_a f_b}$$

- $p_{ab}$ : likelihood of these two residues being correlated because they're homologous.
- $f_a f_b$ : likelihood of these two residues being uncorrelated and unrelated, occurring independently.
- $\lambda$ : a scaling factor

Assuming that each aligned residue pair is statistically independent of the others (biologically dubious, but mathematically convenient):

- The score of an alignment ("**alignment score**") = sum of individual log-odds scores for each aligned residue pair.

# Substitution matrix to measure similarity in sequence alignments

**BLOSUM** (BLOcks SUBstitution Matrix) for protein sequence alignment.

- Scan the BLOCKS database for very conserved regions of protein families (w/o gaps in the alignment).
- Count the relative frequencies of AA and their substitution probabilities.
- Calculate a log-odds score for each of the 210 possible substitution pairs of the 20 standard amino acids.

Ala	4																			
Arg	-1	5																		
Asn	-2	0	6																	
Asp	-2	-2	1	6																
Cys	0	-3	-3	-3	9															
Gln	-1	1	0	0	-3	5														
Glu	-1	0	0	2	-4	2	5													
Gly	0	-2	0	-1	-3	-2	-2	6												
His	-2	0	1	-1	-3	0	0	-2	8											
Ile	-1	-3	-3	-3	-1	-3	-3	-4	-3	4										
Leu	-1	-2	-3	-4	-1	-2	-3	-4	-3	2	4									
Lys	-1	2	0	-1	-3	1	1	-2	-1	-3	-2	5								
Met	-1	-1	-2	-3	-1	0	-2	-3	-2	1	2	-1	5							
Phe	-2	-3	-3	-3	-2	-3	-3	-3	-1	0	0	-3	0	6						
Pro	-1	-2	-2	-1	-3	-1	-1	-2	-2	-3	-3	-1	-2	-4	7					
Ser	1	-1	1	0	-1	0	0	0	-1	-2	-2	0	-1	-2	-1	4				
Thr	0	-1	0	-1	-1	-1	-1	-2	-2	-1	-1	-1	-1	-2	-1	1	5			
Trp	-3	-3	-4	-4	-2	-2	-3	-2	-2	-3	-2	-3	-1	1	-4	-3	-2	11		
Tyr	-2	-2	-2	-3	-2	-1	-2	-3	2	-1	-1	-2	-1	3	-3	-2	-2	2	7	
Val	0	-3	-3	-3	-1	-2	-2	-3	-3	3	1	-2	1	-1	-2	-2	0	-3	-1	4
	Ala	Arg	Asn	Asp	Cys	Gln	Glu	Gly	His	Ile	Leu	Lys	Met	Phe	Pro	Ser	Thr	Trp	Tyr	Val

# Substitution matrix to measure similarity in sequence alignments

**BLOSUM** (BLOcks SUBstitution Matrix) for protein sequence alignment.

Ala	4																			
Arg	-1	5																		
Asn	-2	0	6																	
Asp	-2	-2	1	6																
Cys	0	-3	-3	-3	9															
Gln	-1	1	0	0	-3	5														
Glu	-1	0	0	2	-4	2	5													
Gly	0	-2	0	-1	-3	-2	-2	6												
His	-2	0	1	-1	-3	0	0	-2	8											
Ile	-1	-3	-3	-3	-1	-3	-3	-4	-3	4										
Leu	-1	-2	-3	-4	-1	-2	-3	-4	-3	2	4									
Lys	-1	2	0	-1	-3	1	1	-2	-1	-3	-2	5								
Met	-1	-1	-2	-3	-1	0	-2	-3	-2	1	2	-1	5							
Phe	-2	-3	-3	-3	-2	-3	-3	-3	-1	0	0	-3	0	6						
Pro	-1	-2	-2	-1	-3	-1	-1	-2	-2	-3	-3	-1	-2	-4	7					
Ser	1	-1	1	0	-1	0	0	0	-1	-2	-2	0	-1	-2	-1	4				
Thr	0	-1	0	-1	-1	-1	-1	-2	-2	-1	-1	-1	-1	-2	-1	1	5			
Trp	-3	-3	-4	-4	-2	-2	-3	-2	-2	-3	-2	-3	-1	1	-4	-3	-2	11		
Tyr	-2	-2	-2	-3	-2	-1	-2	-3	2	-1	-1	-2	-1	3	-3	-2	-2	2	7	
Val	0	-3	-3	-3	-1	-2	-2	-3	-3	3	1	-2	1	-1	-2	-2	0	-3	-1	4
	Ala	Arg	Asn	Asp	Cys	Gln	Glu	Gly	His	Ile	Leu	Lys	Met	Phe	Pro	Ser	Thr	Trp	Tyr	Val

- The rarer the amino acid is, the more surprising it would be to see two of them align together by chance.
- L/L more common than WW:
  - $p_{LL} = 0.0371$ ,  $p_{WW} = 0.0065$
- W is a much rarer amino acid:
  - $f_L = 0.099$ ,  $f_W = 0.013$ .

Check with  
 $\lambda = 0.347$ .

# Substitution matrix to measure similarity in sequence alignments

**BLOSUM** (BLOcks SUBstitution Matrix) for protein sequence alignment.

- A/L pairs are slightly more frequent in homologous alignments than K/E pairs:
  - $p_{AL} = 0.0044$ ,  $p_{KE} = 0.0041$ .
- But, A and L are more common amino acids:
  - $p_A = 0.074$ ,  $p_L = 0.099$ ,  $p_K = 0.058$ ,  $p_E = 0.054$ .

Check with  
 $\lambda = 0.347$ .

Ala	4																			
Arg	-1	5																		
Asn	-2	0	6																	
Asp	-2	-2	1	6																
Cys	0	-3	-3	-3	9															
Gln	-1	1	0	0	-3	5														
Glu	-1	0	0	2	-4	2	5													
Gly	0	-2	0	-1	-3	-2	-2	6												
His	-2	0	1	-1	-3	0	0	-2	8											
Ile	-1	-3	-3	-3	-1	-3	-3	-4	-3	4										
Leu	-1	-2	-3	-4	-1	-2	-3	-4	-3	2	4									
Lys	-1	2	0	-1	-3	1	1	-2	-1	-3	-2	5								
Met	-1	-1	-2	-3	-1	0	-2	-3	-2	1	2	-1	5							
Phe	-2	-3	-3	-3	-2	-3	-3	-3	-1	0	0	-3	0	6						
Pro	-1	-2	-2	-1	-3	-1	-1	-2	-2	-3	-3	-1	-2	-4	7					
Ser	1	-1	1	0	-1	0	0	0	-1	-2	-2	0	-1	-2	-1	4				
Thr	0	-1	0	-1	-1	-1	-1	-2	-2	-1	-1	-1	-1	-2	-1	1	5			
Trp	-3	-3	-4	-4	-2	-2	-3	-2	-2	-3	-2	-3	-1	1	-4	-3	-2	11		
Tyr	-2	-2	-2	-3	-2	-1	-2	-3	2	-1	-1	-2	-1	3	-3	-2	-2	2	7	
Val	0	-3	-3	-3	-1	-2	-2	-3	-3	3	1	-2	1	-1	-2	-2	0	-3	-1	4
	Ala	Arg	Asn	Asp	Cys	Gln	Glu	Gly	His	Ile	Leu	Lys	Met	Phe	Pro	Ser	Thr	Trp	Tyr	Val

# Substitution matrix to measure similarity in sequence alignments

Ala	4																			
Arg	-1	5																		
Asn	-2	0	6																	
Asp	-2	-2	1	6																
Cys	0	-3	-3	-3	9															
Gln	-1	1	0	0	-3	5														
Glu	-1	0	0	2	-4	2	5													
Gly	0	-2	0	-1	-3	-2	-2	6												
His	-2	0	1	-1	-3	0	0	-2	8											
Ile	-1	-3	-3	-3	-1	-3	-3	-4	-3	4										
Leu	-1	-2	-3	-4	-1	-2	-3	-4	-3	2	4									
Lys	-1	2	0	-1	-3	1	1	-2	-1	-3	-2	5								
Met	-1	-1	-2	-3	-1	0	-2	-3	-2	1	2	-1	5							
Phe	-2	-3	-3	-3	-2	-3	-3	-3	-1	0	0	-3	0	6						
Pro	-1	-2	-2	-1	-3	-1	-1	-2	-2	-3	-3	-1	-2	-4	7					
Ser	1	-1	1	0	-1	0	0	0	-1	-2	-2	0	-1	-2	-1	4				
Thr	0	-1	0	-1	-1	-1	-1	-2	-2	-1	-1	-1	-1	-2	-1	1	5			
Trp	-3	-3	-4	-4	-2	-2	-3	-2	-2	-3	-2	-3	-1	1	-4	-3	-2	11		
Tyr	-2	-2	-2	-3	-2	-1	-2	-3	2	-1	-1	-2	-1	3	-3	-2	-2	2	7	
Val	0	-3	-3	-3	-1	-2	-2	-3	-3	3	1	-2	1	-1	-2	-2	0	-3	-1	4
	Ala	Arg	Asn	Asp	Cys	Gln	Glu	Gly	His	Ile	Leu	Lys	Met	Phe	Pro	Ser	Thr	Trp	Tyr	Val

**BLOSUM<sub>r</sub>**: the matrix built from blocks with less than **r**% of similarity.

- E.g., BLOSUM62: built using sequences with less than 62% similarity (sequences with  $\geq 62\%$  identity were clustered)
  - BLOSUM62 among the best for detecting most weak protein similarities.
  - Default matrix for protein BLAST.



# Substitution matrix to measure similarity in sequence alignments

Substitution  
matrix for DNA

<b>A</b>				
<b>T</b>				
<b>G</b>				
<b>C</b>				
	<b>A</b>	<b>T</b>	<b>G</b>	<b>C</b>

Making-up an arbitrary matrix by fixing the  $p_{ab}$  values  $\rightarrow$  directly describes what homologous alignments are expected to look like.

- The resulting score matrix is optimal for detecting alignments that match these target frequencies.

Say, the matrix should be optimized for finding 88% identity alignments.

- Assume that all mismatches are equiprobable, and composition of both alignments and background sequences is uniform at 25% for each nucleotide ( $f_a, f_b = 0.25$  for all  $a, b$ ). Then,
  - Four identities:  $p_{aa} = 0.22$
  - 12 types of mismatch:  $p_{ab} = 0.01$ .
- If we set  $\lambda = 1$ , this gives +1.26 for a match and -1.83 for a mismatch.
- Setting  $\lambda = 0.25$  and round off: we have a new scoring system of +4/-7.

# Substitution matrix to measure similarity in sequence alignments

Substitution  
matrix for DNA

A				
T				
G				
C				
	A	T	G	C

Given a scoring matrix, we can back calculate target frequencies if two conditions are met:

$$s(a,b) = \frac{1}{\lambda} \log \frac{p_{ab}}{f_a f_b}$$

1. It must have at least one positive score, and
2. The expected score for random sequence alignments must be negative.

True for most score matrices:

- These properties are necessary to make local sequence alignment algorithms like BLAST and Smith-Waterman work.
- Both conditions are met by definition for matrices derived as log-odds scores, except for the useless case of  $p_{ab} = f_a f_b$  for all  $a, b$ .

Examples:

- FASTA & WU-BLASTN: arbitrary +5/−4 scoring system; Optimal for detecting alignments that are 65% identical.
- NCBI BLASTN: +1/−2 scoring system; Optimal for detecting alignments that are 95% identical.

# How do we scale this up to search an entire sequence database?

Given a query sequence, and a large set of target sequences (millions), which target sequences (if any) are related to the query?

- Individual alignments need not be perfect: Once initial matches are found, they can fine-tune them later.
- Must be very fast.

Exploit the nature of the problem (most sequences will be unrelated to the query):

- If any match with % identity  $\leq 90$  is going to be rejected, can ignore sequences which don't have a stretch of 10 nucleotides in a row.
- Pre-screen sequences for common long stretches.
- Pre-process the database offline and index k-mers.

TITLE

CITED BY

YEAR


## Basic local alignment search tool


SF Altschul, W Gish, W Miller, EW Myers, DJ Lipman

Journal of molecular biology 215 (3), 403-410

136003 \*

1990


 U.S. National Library of Medicine


 National Center for Biotechnology Information

Sign in to NCBI

**BLAST®**
[Home](#)
[Recent Results](#)
[Saved Strategies](#)
[Help](#)

### Basic Local Alignment Search Tool

**BLAST** finds regions of similarity between biological sequences. The program compares nucleotide or protein sequences to sequence databases and calculates the statistical significance.

[Learn more](#)

**NEWS**

**IgBLAST 1.8.0 released**  
 A new version of IgBLAST is now available.  
 Wed, 15 Nov 2017 16:00:00 EST

[More BLAST news...](#)

### Web BLAST



**Nucleotide BLAST**  
nucleotide ► nucleotide

**blastx**  
translated nucleotide ► protein

**tblastn**  
protein ► translated nucleotide



**Protein BLAST**  
protein ► protein

### BLAST Genomes

[Human](#)
[Mouse](#)
[Rat](#)
[Microbes](#)

<https://www.ncbi.nlm.nih.gov/BLAST/>

# Upcoming project deadline: Project topic due on Jan 31

- Discuss with me *and* any other PI; Read recent papers.
- Briefly describe a project idea:
  - Title
  - Project advisor (if someone outside class)
  - 250-word abstract addressing the following 4 Qs:
    - What is the problem?
    - How is it addressed currently & what are the limitations?
    - What is your approach to addressing it & why is likely to be successful?
    - If successful, why does it matter (what is the impact)?
- Submit a PDF.